Master Thesis: Open Specification of a user-controlled Web Service for Personal Data

G. Jahn

November 10, 2016

Abstract

Data is the currency of tomorrow. Organizations, whether in the private
or public sector, gathering enormous amounts of personal (big) data. This
data is harvested and incorporated by these third parties, but were created
by individuals and therefore should belong to them. People depending on
all their data. Their identity as well as their personality are defined by their
personal data. Meanwhile data silo operators are hammering onto these
haystacks eager trying to find any kind of correlations worth interpreting,
thereby almost inevitably discriminating the rightful owners. To reduce the
possibility of discrimination only the least amount of data that is required
should be handed over to the third party. Thus the individual has to be in
charge of the whole process. A personal data service will empower its user
to regain full control over her data and facilitates detailed information on
every data flow. To be able to trust such a tool, the user should be able
to look inside. Therefore a personal data service has to be open source and
developed transparently, which then also allows to self-host it.

Contents

1

Introduction

1.1 Motivation

Nowadays it is rare to find someone who doesn't collects data about some kind of things, particularly humans are the targets of choice for the *Big Data Movement.* But since humans are all individuals, which means we are all distinct from each other - more or less. Some of us are sharing a bit more similarities, but most are much less similar to each other. Since these few similarities are widely shared, they should be less important, because it seems to be intuitively the nature of inflationary occurrence, but instead

the opposite happens to be the case. It allows to determine who is part of a subset with a specific, and therefore shared, attribute and who isn't, in order to relate apparently common known stereo typical patters onto the individuals in that subset just to predict outcomes of the corresponding problem or question. In other words, searching for causation where in best case one might find correlations - or so called *discrimination*, which

> [...] refers to unfair or unequal treatment of people based on membership to a category or a minority, without regard to individual merit. *[1]*

When interacting directly with each other, discrimination of human beings is still a serious issue in our society, but also when humans leveraging computers and algorithms to uncover former unnoticed information in order to include them in their decision making. For example when granting credits, hiring employees, investigating crimes or renting flats. Approving or denying is all happening based on data about the affected individuals [2], which is nothing but discrimination, but on a much larger scale and with less effort - almost parenthetically. The described phenomenon is originally referred as *Bias in computer systems* [3]. What at fist seems to look like machines going rouge on humans, is in fact the *cognitive bias* [4] of the human nature, modeled in machine executable language and made to reveal the patterns their creators were looking for.

In addition to the identity-defining data mentioned above humans have the habit to create more and more data on a daily basis - pro-actively e.g by writing a tweet and passively e.g by allowing the twitter app accessing their current location while submitting the tweet. As a result already tremendous

5

amounts of data keeps growing bigger and bigger waiting to get harvested, collected, aggregated, analyzed and finally interpreted. The crux here is, the more data being made available [5] to *mine* on, the chances are much higher to isolate data sets, that differ from each other but are coherent in theirselves. Then it is just a matter of the preceded questioning how to distinguish the data set and thereby the related individuals from each other.

So to lower potential discrimination we either need to erase responsible parts from the machines, therefore raising awareness and teaching people about the issue of discrimination are crucial, or we*re trying to prevent our data from falling into these data silos. The later will be addressed in this work.

1.2   Purpose & Outcome

From an individual's perspective providing data to third parties might not seem harmful at all. Instead they eventually get improved services in return, e.g. more adequate recommendations and fitting advertisement, or more helpful therapies and more secure environments. That said, though it is a matter of perception what's good and bad, what's harmful and what's an advantage. Computing data to leverage decision making is essentially just science and technology and it's up to the humans how such tools are getting utilized and what purposes they are serving. Hence it should be decided by the data creators, how their data get processed and what parts of them are used.

To tackle the described issue the initial idea here is (1) to enable individuals with the ability to control and maintain their entire data distribution and (2) thus reducing the amount of *potentially discriminatory* [1] attributes leaking

into arbitrary calculations. To do so people need a reliable and trustworthy tool, which helps to manage all their *personal data*

respect privacy and is secure

The way how to solve the described dilemma is not new. Early days of work done in this field can be dated back to the Mid-2000s where studies were made e.g. about recent developments in the industry or user's concerning about privacy, and the term *Vendor Relationship Management (VRM)* were used initially within the context of user-centric personal data management, which also led into the *ProjectVRM* [6] started by the *Berkman Klein Center for Internet & Society at Hardvard University.* Since then a great amount of effort went into this research area until today, while also commercial products and business models trying to solve certain problems. For instance concepts such as the *Personal Data Store (PDS)* [7] or a *MyData* [8] implementation called *Meeco* [9], which will all be covered in a more detailed way within the following chapter.

The work and research done for this thesis is intended to be the foundation for an *Open Specification* which by itself is a manual to implement a *Personal Data as a Service.*

Furthermore to ensure a minimum level of trust

it is vital to make the development process fully transparent, encourage people to get involved and of every

- Open Source

The approach discussed in this document shall be the initial step in creating a tool to manager all the data that defines her identity controlled and ad-

ministrated by it's owner, and maybe give her a more precise understanding about where her personal information flows and how this might effect her privacy.

1.3   Terminologies

- serverless: https://auth0.com/blog/2016/06/09/what-is-serverless/
- digital footprints: TODO
- owner: person who controls (and probably hosts) the data service containing her personal data
- (data) consumer: Third party, external entity requesting data, authorized by the owner to do so

2

Fundamentals

2.1 Personal Data in the context of the Big Data Movement

- difference between users *profile/account data* and their *meta data*?

2.2 Personal Data as a Product

...to do what exactly? (NOTE: was previously in *Motivation*)

- If one would ask collectors for their motivation to do so, most likely the answer would be something along the lines of *"We want to have a better understanding of our customers"*. To do what exactly?

- To predict what might be the next thing I am supposed to buy Or what things I probably would like to consume but most certainly not yet know of?

- Let's take a look at some examples. An advertising service uses tracking data for targeted advertising. The more data they collect about an individual, the more accurate decisions they are able to make on what ads they have to display are the ones the individual will click on and then will go all the way down to disclose a successful purchase. This makes the placed advertisement more valuable for google and therefore more expensive to the advertiser, because of a high precision. Or a streaming provider content recommendation is also based on heavy user profiling done by looking at her consumption history, tracked platform interactions and probably many more vectors. As always, it depends on the business model, but it seems to be consensual, that it all comes down to improving and enhancing the collector's product (NOTE: needs at least some empirical evidence), in order to satisfy the customers - and that on the other hand depends on who is seen as customers.

- Nevertheless if we change the perspective, a lot of great things can be achieved with the help of huge amounts of personal data, such as:

  - planing and managing human resources for situations, like e.g. big events where attendees might need some help [10]

- predicting infrastructure workloads [http://ieeexplore.ieee.org/document/7336197/]

- making more accurate diagnostics to improve their therapy [11]

- finding patters in climate changes, which otherwise wont get revealed [12].

- individuals then get in role of selling/offering it's own data to those who were previously collecting them

- It is all about understanding the human being and why she behaves as she does. The challenge is not only to computing certain motives but rather concluding to the right ones. When analyzing computed results with the corresponding data models and trying to conclude, it is important to keep in mind, that correlation is no proof of causation.

2.3 Personal Data as of the Law

- who is the owner in what situation or under what circumstance?
- difference between "Personal Data" and the owner of that data?

2.4 Digital Identity

- identity defining data (e.g. history of personal ID card)

- with such a system a human being is represented by a non-physical abstraction of herself. Which essentially is a list of attributes, that are at least for legal and civil administration purposes important. Their values in total are unique and representing the corresponding human.

Certain attributes hold unique values within it's own context, for example the *social security number.*

- Thus it's not necessary to know the values of all attributes in order to identify it's owner

- therefore its imported to not see it as a reduction of a living individual to some bits and bytes

- what will happen with her data service after a person died?

2.5 Use Cases (NOTE: maybe move to 100_introduction)

- package shipment after buying sth online
- social network accessing arbitrary profile data
- making an online purchase
- credibility (requesting credit permission) validation by a certain financial institution: accessing arbitrary data
- patient/health record
- care (movement) data

2.6 Related Work

- much more research since the data mining, big data, deep learning

2.6.1 RESEARCH

- openPDS/safeAnswer [http://openpds.media.mit.edu/]
- TAS3 aka ZXID aka Synergetics (lead arch Sampo Kellomäki also Co-Authored openPDS papers)
- Higgins [https://www.eclipse.org/higgins/]
- Hub-of-All-Things [http://hubofallthings.com/what-is-the-hat/]
- ownyourinfo [http://www.ownyourinfo.com]
- PAGORA [http://www.paoga.com]
- PRIME/PrimeLife [https://www.prime-project.eu, http://primelife.ercim.eu/]
- databox.me (reference implementation w/ the "solid" framework)
- Microsoft HealthVault
- Industrial Data Space (german research project mainly driven by Fraunhofer Institute)
- Polis (greek research project from 2008) [http://polis.ee.duth.gr/Polis/index.php]

2.6.2 ORGANISATIONS

- Kantara Initiative (former "Liberty Alliance") [https://kantarainitiative.org/]
- Open Identity Exchange [http://openidentityexchange.org/resources/white-papers/]
- Qiy Foundation [https://www.qiyfoundation.org/]

2.6.3 COMMERCIAL PRODUCTS

- MyData [https://mydatafi.wordpress.com/]
- Meeco (killing the ad provider middle man) [https://meeco.me/how-it-

works.html]

- RESPECT network [https://www.respectnetwork.com/]

- aWise AEGIS [http://www.ewise.com/aegis]

2.7   Standards and Specifications

- http(s)

- all the *Semantic Web* stuff

- Container/App spec

- JWT

- oAuth (?)

- JSON

- REST

- GraphQL

3

Core Principles

NOTE: here we discuss a variety of possibilities –> conceptual work

3.1   Data Ownership

- user-centric, full control

3.2 Identity Verification

- maybe go with a Signing/verifying Authority (aka CA)
  - do I trust the gov or certain companies more? Which interests do these Role/Stakeholder have?
  - revoking the cert which provides the authenticity of the individual's digital identity should only be possible with a two-factor secret. One part of this secret is owned by the CA and the other half has the individual behind the personal API
- TODO: look into
  - PKI
  - ePerso
  - E-Post/de-mail
- Authentication

3.3 Authorisation (remove in favour of data access?)

- NOTE: does not mean this tool authenticates it's owner against third party platforms like OpenID does. but it could play the role of the 2n factor in a multi-factor authentication process (if the mobile-device-architecture was chosen)
- refers primarily to the process of a data consumer (third party, which needs the data for whatever reason) verifies her admission to request

3.4 Authentic Data

- is this data (in this case identity) certified or not (results in higher value)

3.5 Supervised Data Access

- pure/plain data request/resonse
- remote computation/execution (assuming there is no client for the consumer) like https://webtask.io/

3.6 Encapsulation

- containerization (coreos, rkt, mirageos aka unikernal)

3.7 Open Development

- which and why open standards
- why open source
- collaborative transparent development
- Hosting & Administration
  - DYI
  - Usability

4

Requirements

4.1   User Interaction

- as effortless as possible

4.2   Management & Organisation

4.3   Administration

5

Design

5.1   Architecture

- showing possible directions, e.g.:
    - cloud or local storage
    - which components can go where
    - remote execution, to prevent data from leaving the system

5.1.1 OVERVIEW

- distributed architecture (e.g. notification/queue server + mobile device for persistence and administration)

5.1.2 COMPONENTS

5.1.3 PLUGINS

- but for what? and not harm security at the same time? maybe just for data types and admin UI to display analytical (time based) data in other ways
- what about extensions (see iOS 10) to let other apps consume data; only on a mobile device and only if the data is stored there

5.2 Data

- keep in mind to make it all somehow extendible, e.g. by using and storing corresponding schemas
- NOTE: step numbers marked with a * are somehow tasks which are happening in the background and don't require any user interaction

5.2.1 Modelling

5.2.2 Categories (or Classes)

5.2.3 Types

5.2.4 Persistence

- database requirements

5.2.5 Access & Permission

- data needs to have an expiration date

IF01 - Authorizing a consumer to request certain data

1) owner creates a new endpoint URI (like *pdaas.ownersdomain.tld/e/consumer-name*) within the *management user interface*

2) owner passes this URI on to the *consumer*, e.g. through submitting a form or using any arbitrary, eventually insecure channel 3*) consumer need to call this URI for the fist time to verify it's authenticity

3) owner then get's a notification which asks her for permissions to access certain data under the listed conditions 5*) consumer will be informed about the outcome of the owner's decision (NOTE: alongside with some details? how do they look like? XXX need to be in the spec)

5.2.6 CONSUMPTION (DATA INFLOW)

- how data will get into the system

- hwo is the user able to do that, and how does it works

5.2.6.1 Manually

5.2.6.2 Automatically

5.2.7 EMISSION (DATA OUTFLOW)

- depending on what category of data, they might need to get anonymized somehow before they leave the system

- oAuth (1.0a and 2) requires consumers to register upfront. Since the current flow indicates that the initial step is done by the owner, that would cause an overhead in user interactions. Although the owner already *authorized* the consumer simply by submitting a unique URI (`pdaas-server.tld/register?crt=CONSUMER_REGISTER_TOKEN`), of which the `crt` is considered private. Even though the registration provides the consumer with mandatory information such as a consumer identifier (v1: `oauth_consumer_key`, v2: `client_id`) and, depending on the client type, a secret (see https://tools.ietf.org/html/rfc6749#section-2), this process it is not part the specification (https://oauth.net/core/1.0a/#rfc.section.4.2, https://tools.ietf.org/html/rfc6749#section-2). This enables the possibility of integrating oAuth into the consumer registration flow

by using the `CONSUMER_REGISTRATION_TOKEN` as oAuth's *client identifier*. The lack of credentials (v1: `auth_consumer_secret`, v2: `client_secret`) would require transferring the consumer identifier done over a secure channel (e.g. TLS). That would leave *oAuth2* as the version of choice, since it relies on *HTTPS* adn therefore makes the *secret* optional. Where on the other side oAuth 1.0a requires a *secret* to create a signature in order to support insecure connections..

- A general and URI for 3rd parties to register (aka requesting authentication) would raise the issue of dealing with spam request and how to distinct these from the actual ones.

5.2.8 HISTORY

- data versioning
- access logs

5.3 Interfaces

5.3.1 INTERNAL

- UI for Management & Administration

5.3.2 EXTERNAL

- should there be a way to somehow request information about what data is available/queryable, or would this be result in spam/crawler and security issues (also a question for the topic of permissions/sensibility level of certain data)

- certain types of requests, depending on expire date:

  - "ask me any time"
  - "allowed until further notice"
  - one-time permission (but respecting certain http error codes and possible timeout - that might not count)

6

Specification

- what does *open* in Open Specification even mean?

6.1 Processes (TODO: find another word; "Protocol flows"?)

6.2 Application Programming Interfaces

6.3 Graphical User Interfaces

6.4 Security

- the downside of having not just parts of the personal data in different places (which is currently the common way to store), is in case of security breach, it would increase the possible damage by an exponential rate Thereby all data is exposed at once, instead of not just the parts which a single service has stored

- does it matter from what origin the data request was made? how to check that? is the requester's server domain in the http header? eventually there is no way to check that, so me might need to go with request logging and trying to detect abnormal behaviour

- is the consumer able to call the access request URI repeatedly and any time? (meaning will this be stateless or stateful?)

- initial consumer registration would be done on a common and valid https:443 CA-certified connection. after transferring their cert to them as a response, all subsequent calls

need to go to their own endpoint, defined as subdomains like
`consumer-name.owners-notification-server.tld`

6.4.1 Environment

6.4.2 Transport

- communication between internal components *must* be done in https only, but which ciphers? eventually even http/2?

6.4.3 Storage

6.4.4 Authentication

- how should consumer authenticate?

6.5 Recommendations

6.5.1 Software Dependencies

6.5.2 Host Environment

7

Conclusion

7.1 Ethical & Social Impact (TODO: or "Relevance")

7.2 Business Models & Monetisation

- possible resulting direct or indirect business models
- owner might want to sell her data, only under her conditions. therefore some kind of infrastructure and process is required (such as payment transfer, data anonymization, market place to offer data)

7.3 Challenges

7.4 Solutions

7.5 Attack Scenarios

7.6 Future Work

- maybe enable the tool to play the role of an own OpenID provider?

7.7 Summary

- main focus
- unique features
- technology stack & standards
- resources
- the tool might be not a bulletproof vest, but

[1] D. Pedreshi, S. Ruggieri, and F. Turini, "Discrimination-aware data mining," in *Proceedings of the 14th ACM SIGKDD international conference on Knowledge discovery and data mining*, 2008, pp. 560–568 [Online]. Available: http://dl.acm.org/citation.cfm?id=1401959. [Accessed: 03-Nov-2016]

[2] S. Spiekermann, *Ethical IT Innovation: A Value-Based System Design Approach.* CRC Press; Taylor & Francis Group, LLC, 2015, pp. 66–72 [Online]. Available: https://www.crcpress.com/

Ethical-IT-Innovation-A-Value-Based-System-Design-Approach/Spiekermann/
p/book/9781482226355

[3] B. Friedman and H. Nissenbaum, "Bias in computer systems," *ACM Transactions on Information Systems (TOIS)*, vol. 14, no. 3, pp. 330–347, 1996 [Online]. Available: http://dl.acm.org/citation.cfm?id=230561. [Accessed: 07-Nov-2016]

[4] "Cognitive bias," *Wikipedia*, Oct-2016. [Online]. Available: https://en.wikipedia.org/w/index.php?title=Cognitive_bias&oldid=742803386. [Accessed: 08-Nov-2016]

[5] A. Dewes, "C3TV - Say hi to your new boss: How algorithms might soon control our lives." 29-Dec-2015. [Online]. Available: https://media.ccc.de/v/32c3-7482-say_hi_to_your_new_boss_how_algorithms_might_soon_control_our_lives#video&t=1538. [Accessed: 03-Nov-2016]

[6] "ProjectVRM - about. ProjectVRM," 25-Feb-2010. [Online]. Available: https://blogs.harvard.edu/vrm/about/. [Accessed: 09-Nov-2016]

[7] Tom Kirkham, Sandra Winfield, Serge Ravet, and S. Kellomaki, "The personal data store approach to personal data security," *IEEE Security & Privacy*, vol. 11, no. 5, pp. 12–19, 2013.

[8] A. Poikola, K. Kuikkaniemi, and H. Honko, "MyData – a nordic model for human-centered personal data management and processing," pp. 1–12, 2014 [Online]. Available: https://www.lvm.fi/documents/20181/859937/MyData-nordic-model/2e9b4eb0-68d7-463b-9460-821493449a63. [Accessed:

10-Nov-2016]

[9] "Meeco how it works." [Online]. Available: https://meeco.me/how-it-works.html. [Accessed: 09-Nov-2016]

[10] J. Ao, P. Zhang, and Y. Cao, "Estimating the Locations of Emergency Events from Twitter Streams," *Procedia Computer Science*, vol. 31, pp. 731–739, 2014 [Online]. Available: http://linkinghub.elsevier.com/retrieve/pii/S1877050914004980. [Accessed: 05-Nov-2016]

[11] G. Palem, "The Practice of Predictive Analytics in Healthcare," *ResearchGate*, Apr. 2013 [Online]. Available: https://www.researchgate.net/publication/236336250_The_Practice_of_Predictive_Analytics_in_Healthcare. [Accessed: 05-Nov-2016]

[12] N. Burger, B. Ghosh-Dastidar, A. Grant, G. Joseph, T. Ruder, O. Tchakeva, and Q. Wodon, "Data Collection for the Study on Climate Change and Migration in the MENA Region," 2014 [Online]. Available: https://mpra.ub.uni-muenchen.de/56929/. [Accessed: 04-Nov-2016]